

Chess Assignment

van den Bergh Laurin

Institute of Informatics

University of Zurich

Zurich, Switzerland

laurin.vandenbergh@uzh.ch — Matriculation number: 16-744-401

Abstract—This document is a model and instructions for L^AT_EX. This and the IEEEtran.cls file define the components of your paper [title, text, heads, etc.]. *CRITICAL: Do Not Use Symbols, Special Characters, Footnotes, or Math in Paper Title or Abstract. 1 sentence summary, 1 sentence method, and 1 sentence results, and 1 sentence conclusion.

Index Terms—reinforcement learning, chess

I. INTRODUCTION

A. - describe general setting

In this assignment, we explore three different reinforcement learning algorithms to learn how to play a simplified version of chess. These three algorithms are SARSA, Q-Learning and DQN¹.

B. - touch on methods

Throughout the report we indicate in footnotes which task a particular section is referring to in terms of answering the task. We do this as the solutions to certain tasks are spread throughout multiple sections, e.g. task 3 is answered in Section II and Section III. Even though this assignment was not solved in a group, we decided to also answer some of the “group only” and we stick to the numbering of the assignment in order to avoid confusion.

C. - touch on results and conclusion

seeds and nonseeded experiments foreshadow results from qlearning

II. METHODS

A. Environment

This version of chess takes place on a 4 by 4 board and can be thought of as a late game where the agent has a king and a queen, and the opponent has only a king. Since this game can only end in a win for the agent or in a draw, it is the agent’s goal to learn how to win the game and avoid draws. For all experiments considered, the agent will be given a reward of 1 for winning, 0 for drawing, and 0 for all intermediate steps.

For DQN no sequence is needed, because chess fulfills Markov property and gives rise to a Markov Decision Process (MDP). Also no preprocessing of sequence is needed anymore because state encoding is already given. [1] [2]

Chess has/fulfills the Markov property, so applying reinforcement learning algorithms like Q-Learning and SARSA is theoretically justified.

(1) Describe the algorithms Q-learning and SARSA. Explain the differences and the possible advantages/disadvantages between the two methods. No coding is needed to reply to this question.

B. SARSA and Q-Learning²

1) Temporal-Difference Algorithms: maybe shorten SARSA and Q-Learning are two very related model-free types of temporal-difference (TD) algorithms for learning expected rewards, also known as Q-values, when rewards are not immediate and possibly sparse. The learning takes place via interaction with an environment through trial and error. These Q-values are in general represented by an action-value function Q and, for finitely many state-action pairs (s, a) , can be considered as a Q-table where each state-action pair, (s, a) , maps to a single Q-value, thus providing an estimate of the quality of any given state-action pair (s, a) . In this assignment however we use neural networks to approximate the action-value function, which outputs the Q-values for all possible actions for any given state. This helps to avoid computing large Q-tables. All algorithms explored in this assignment, including DQN, require the environment to fulfill the Markov property. why? else what?

2) differences/what they have in common: SARSA and Q-Learning address the temporal-credit-assignment problem [3], that is, trying to attribute future rewards to previous actions. These future rewards get discounted with the hyper-parameter γ (see Section III-C). Both algorithms repeatedly choose and take actions in the environment according to some policy π , e.g. an ϵ -greedy policy.

However, this is where they differ. SARSA is an on-policy algorithm, which means that consecutive actions are chosen according to the same policy π , even during the update step of the Q-values, which leads to the update rule:

$$Q_{\pi}(s, a) \leftarrow Q_{\pi}(s, a) + \eta(r + \gamma Q_{\pi}(s_{t+1}, a') - Q_{\pi}(s, a))$$

for some future action a' chosen according to policy π .

Q-learning, on the other hand, is an off-policy algorithm, which means that it takes its actions a according to its policy π , but during the update steps it assumes a greedy policy, i.e.

¹SARSA serves as answer to task 3 and DQN serves as answer to task 5.

²Answer to task 1.

optimal play, for future actions a' . Q-Learning has the update rule:

$$Q_\pi(s, a) \leftarrow Q_\pi(s, a) + \eta(r + \gamma \max_{a'} Q_\pi(s_{t+1}, a') - Q_\pi(s, a)).$$

3) - *advantages/disadvantages*: This leads to one of Q-Learning's major advantages: Because of Bellman's optimality equation, Q-Learning is guaranteed to learn the values for the optimal policy, i.e. $Q_*(s, a) = \max_\pi Q_\pi(s, a)$, and in a greedy setting will take the optimal actions, at least if it was trained sufficiently. However, this can in certain cases mean that the online performance of Q-Learning will be worse than the one from SARSA, as Sutton et al. [4] demonstrate with their "gridworld" example "Cliff Walking". Our chess game is a similar situation, because a win and a draw can be very close, thus during exploration Q-Learning can accidentally create a draw because it is going for the optimum when exploiting. Q-Learning is however relatively unstable and the parameters can even diverge when it is combined with non-linear function approximators [2], making the guarantee to learn the optimal policy irrelevant.

SARSA will learn to take a safer path, because it keeps its policy in mind when updating the Q-values, i.e. it keeps in mind that it will explore in future actions. This has the advantage that SARSA in general tends to explore more than Q-Learning.

C. Experience Replay³

replay memory of fixed size, queue (get code in appendix)

(2) [Group Only] Describe the experience replay technique. Cite relevant papers.

Experience replay is a technique proposed by Lin [5] to speed up the training process for reinforcement learning algorithms by reusing past experiences for future training. This is analogous to the human ability to remember past experiences and learn from them even after the fact. The past experiences are stored in a replay memory of fixed size at each time step t as a tuple $e_t = (s_t, a_t, r_t, s_{t+1})$. This essentially allows us to transform the learning process from online learning to mini-batch learning, where a batch of experiences e_j is randomly sampled for each update step. Experience replay can only be used in combination with off-policy algorithms, because otherwise the current parameters determine the next sample and create unwanted feedback loops [1], [2].

Experience replay provides many benefits over online Q-Learning, especially when neural networks are used to approximate the action-value function. First, it enables the agent to learn from past experiences more than once, leading to increased data efficiency and faster convergence [1], [5]. Second, since for each update step past experiences are sampled randomly, the correlations between the individual actions are reduced, which then reduces the variance of the updates [1]. This leads to the experience samples e_j being closer to i.i.d. and thus guaranteeing better convergence when using optimization algorithms such as stochastic gradient descent as most convergence proofs assume i.i.d. data.

³Answer to "group only" task 2.

D. Deep Q-Networks (DQN)⁴

A first version of the DQN algorithm was proposed by Mnih et al. [1] and combined experience replay with Q-learning, where a neural network was used as a non-linear function approximator for the action-value function. Mnih et al. [2] later improved upon the method and presented the DQN algorithm, as it is known today, where they address the problem of the Q-values $Q(s, a)$ being correlated to the target values $y = r + \gamma \max_{a'} Q(s', a')$ because they are generated using the same neural network. In the DQN algorithm they separated the Q-network from the target network and only update the target network every C steps, which helps to break this correlation and combat diverging network parameters.

E. Experiments

In order to address all tasks, we divided the tasks into several independent experiments. First, we conducted seeded runs⁵ for all three algorithms using seed 21 for reproducibility, which was chosen a-priori. These seeded runs serve as examples to compare the algorithm's online performance qualitatively. The seeds are used such that the weights of all neural networks are instantiated identically for all algorithms and they subsequently serve as seeds for any random number used during training. This makes sure that all agents start with the same initial conditions and that the results are reproducible (see Section V-1). All algorithms were run for 100000 episodes using identical model architecture and hyper-parameters (see Section II-F).

Since the seeded runs are heavily influenced by the choice of the seed, we could end up with anything between a very lucky and well performing seed, or with a very unlucky one. Also the interpretation of the seeded runs is more difficult as we just have one run for each algorithm. Therefore, we decided to perform a simulation study and complete 30 non-seeded runs for each algorithm in order to get a better idea of how the algorithms perform on average. For computational reasons we limited these runs to 40000 episodes as we realized with test runs that by then most of the training progress has already taken place.

To analyze the impact of the hyper-parameters β and γ ⁶ we trained 49 agents with different combinations for β and γ but keeping all other hyper-parameters and model architecture identical. We chose SARSA for this experiment as we found it to have very low variance between its unseeded runs, which makes it an ideal candidate for comparing individual runs (see Figures 4 and 5). These runs are seeded identically to the seeded runs mentioned above.

F. Implementation and Hyper-parameters

1) *Implementation*: We implemented all algorithms from scratch according to Sutton et al. [4] (SARSA and Q-

⁴Answer to task 5: Describing the used method.

⁵Answers to task 3 and 5.

⁶Answer to task 4.

Learning⁷) and Mnih et al. [2] (DQN⁸). For the implementation see file `neural_net.py` on GitHub or Listing 1. All algorithms use a neural network with 58 input neurons, 200 hidden neurons and 32 output neurons, not including the biases for the input and hidden layer. The neural network automatically adds a constant input for the bias and the hidden layer. The implementation treats the biases like any other weights and thus they are part of any matrix multiplication. We used a ReLU activation function for the hidden layer and no activation on the output layer. The weights were initialized using Glorot initialization [6], such that the weights are sampled from a normal distribution with mean 0 and variance $\frac{2}{n_{in}+n_{out}}$, where n_{in} and n_{out} denote to the number of input and output neurons of the respective layer. This helped preventing exploding gradients for the most part.
check again with results

2) *Hyper-Parameters*: For all experiments we used the default hyperparameters provided in the `Assignment.ipynb` file unless otherwise noted. For DQN we updated the weights of the target network after every $C = 10$ steps, as most games take fewer steps than that. We used a replay memory of size 100000 and a batch size of 32.

III. RESULTS

A. Seeded Runs⁹

(3) We provide you with a template code of the chess game. Fill the gaps in the program file `chess` implementing either Q Learning or SARSA; detailed instructions are available inside the file. We provide indicative parameter values. Produce two plots that show the reward per game and the number of moves per game vs training time. The plots will be noisy. Use an exponential moving average.

The rewards and number of moves for the seeded runs are depicted in Figures 1 and 2 respectively. Since the curves are very noisy, we smoothed them using an exponential moving average (EMA) with a weight on the most recent observation of $\alpha = 0.001$.

As expected, the online performance of Q-Learning in terms of the rewards is generally lower than the rewards for SARSA but they converge slowly as ϵ decreases (Figure 1). Also in Figures 1 and 2 we can see that Q-Learning experiences instable learning behavior as both plots are a lot more noisy and at about 20000 and 90000 episodes the rewards decrease for some period. SARSA and DQN don't show this behavior.

Even though the number of steps is not punished, all agents still learn to reduce the number of steps over time, as they do not give rewards and their goal is to take actions that do. SARSA seems to do the best job at this, which perhaps is caused by its tendency to explore more and find better strategies. Q-Learning however seems to struggle to reduce the number of steps it takes.

⁷SARSA as answer to task 3 and Q-Learning as additional algorithm.

⁸Answer to task 5.

⁹Answer to task 3 (SARSA and Q-Learning as additional algorithm) and 5 (DQN).




Fig. 1: Exponential moving average of the rewards achieved during training. All three algorithms, Q-Learning, SARSA and DQN were initialized with identical weights and trained with identical network architecture and hyper-parameters.




Fig. 2: caption

(5) Implement another Deep RL algorithm (it could be SARSA if you chose before Q Learning and vice versa). Compare the new results with your previous results.

As suggested by Mnih et al. [1], [2], DQN¹⁰ was able to overcome the downsides of Q-Learning which lead to an online performance which is comparable to that of SARSA in terms of reward and number of moves it achieved, and also in terms of the stability during training. It did however not learn to reduce the amounts of steps as much as SARSA.

B. Simulation Study (Non-seeded Runs)

We can confirm that the qualitative results from the seeded runs reasonably well represent the average case. The only notable exception being Q-Learning, for which most runs performed equally well to the seeded run, but some runs experienced huge increases in the number of steps, which influenced the average run dramatically, leading to an average of about 23 moves per episode after 40000 episodes. For the learning curves see Figures 4 and 5 in the Appendix.

We observed that DQN and SARSA show very comparable learning curves with DQN showing slightly faster convergence in the first 5000 episodes. SARSA showed the lowest variance in all runs and seems to be a very stable algorithm. Q-Learning on the other hand showed clear signs of divergence as for

¹⁰Answer to task 5 for comparing DQN to SARSA and Q-Learning.

some runs the rewards consistently dropped while the number of moves consistently increased. This shows that the measures taken by Mnih et al. [2] to combat the disadvantages of Q-Learning worked and increased the stability as well as the convergence speed.

We can conclude that the seeded runs in our initial experiment truthfully represent the average run and therefore some level of inference is justified.

We also found out that, unsurprisingly, the effective training time is mainly dependent on the number of steps an algorithm takes per episode. This leads to Q-Learning having by far the longest training time, especially when the parameters diverge and the number of steps increase. DQN and SARSA have relatively short training times, with SARSA being the fastest.

C. Hyper-parameters

(4) Change the discount factor γ and the speed β of the decaying trend of ϵ (for a definition of β please see code). Analyse the results.

1) *gamma*: Figure 3(a) depicts the rewards and number of moves per episode as a function of β and γ . We can see that the reward increases monotonically as γ is increases, suggesting that a value of $\gamma \in [0.80, 1)$ should be chosen. This intuitively makes sense, as we have very sparse rewards and want the agent to “backpropagate” this reward through its sequence of actions.

2) *beta*: We can not see a clear relationship between β and the rewards, apart from $\beta = 0$ being an inferior choice. In Figure 3(b) we can see that the that the number of steps taken by the agent decreases drastically when increasing γ from very low levels, but again there is no clear pattern visible for β . In summary, for reasonalby chosen values of γ the choice of β seems to not have much of an influence for training periods up to 40000 episodes.

(6) [Group Only] Change the state representation or the administration of reward. Interpret your results.

(7) [Group Only] The values of the weights could explode during learning. This issue is called exploding gradients. Explain one possible way to fix this issue and implement it. For example, search and study RMSprop). Show in a plot how your solution fixed this issue.

IV. CONCLUSION

We are aware that the performance of the individual algorithms could be improved by tuning the hyper-parameters, however, this was not explicitly asked for and the focus on this assignment lies on the comparison of these algorithms.


summarize results, takeaway

get correct bib format and complete entries

REFERENCES

- [1] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, and M. A. Riedmiller, “Playing atari with deep reinforcement learning,” *CoRR*, vol. abs/1312.5602, 2013. [Online]5 Available: <http://arxiv.org/abs/1312.5602>

1A Reward for different values of beta and gamma (SARSA, 40000 episode)



(a)

umber of Moves for different values of beta and gamma (SARSA, 40000 e




Fig. 3


- [2] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. A. Riedmiller, A. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik, I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis, “Human-level control through deep reinforcement learning,” *Nat.*, vol. 518, no. 7540, pp. 529–533, 2015. [Online]. Available: <https://doi.org/10.1038/nature14236>
- [3] R. S. Sutton, “Temporal credit assignment in reinforcement learning,” 1984.
- [4] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. The MIT Press, 2018. [Online]. Available: <http://incompleteideas.net/book/the-book-2nd.html>
- [5] L.-J. Lin, “Self-improving reactive agents based on reinforcement learning, planning and teaching,” *Mach. Learn.*, vol. 8, no. 3–4, p. 293–321, may 1992. [Online]. Available: <https://doi.org/10.1007/BF00992699>
- [6] X. Glorot and Y. Bengio, “Understanding the difficulty of training deep feedforward neural networks,” in *In Proceedings of the International Conference on Artificial Intelligence and Statistics (AISTATS’10)*. Society for Artificial Intelligence and Statistics, 2010.

IEEE conference templates contain guidance text for composing and formatting conference papers. Please ensure that all template text is removed from your conference paper prior to submission to the conference. Failure to remove the template text from your paper may result in your paper not being published.


V. APPENDIX

1) *Reproducibility*: put the appendix after the bibliography?


```
# import libraries
from types import MethodDescriptorType
import numpy as np
from tqdm.notebook import tqdm
import os
import json
```

(a)




(b)




(c)


Fig. 4



(a)



(b)



(c)

Fig. 5

```

7 import time
8 import random
9 from collections import namedtuple, deque
10
11 # import from files
12 from Chess_env import *
13
14
15
16 # ===== Epsilon-greedy Policy =====
17
18 def EpsilonGreedy_Policy(Qvalues, allowed_a, epsilon
19 ):
20     """
21     returns: tuple
22         an action in form of a one-hot encoded
23         vector with the same shape as
24         Qvalues.
25         an action as decimal integer (0-based)
26
27     Assumes only a single state, i.e. online
28     learning and NOT (mini-)batch learning.
29     """
30     # get the Qvalues and the indices (relative to
31     # all Qvalues) for the allowed actions

```

```

27 allowed_a_ind = np.where(allowed_a==1)[0]
28 Qvalues_allowed = Qvalues[allowed_a_ind]
29
30
31 # ----- epsilon greedy -----
32
33 # draw a random number and compare it to epsilon
34 rand_value = np.random.uniform(0, 1, 1)
35
36 if rand_value < epsilon: # if the random number
37     is smaller than epsilon, draw a random
38     action
39     action_taken_ind_of_allwed_only = np.random.
40         randint(0, len(allowed_a_ind))
41 else: # greedy action
42     action_taken_ind_of_allwed_only = np.argmax(
43         Qvalues_allowed)
44
45 # get index of the action that was chosen (
46     relative to all actions, not only allowed)
47 ind_of_action_taken = allowed_a_ind[
48     action_taken_ind_of_allwed_only]
49
50 # ----- create usable output -----

```

```

46 # get the shapeensions of the Qvalues 115
47 N_a, N_samples = np.shape(Qvalues) # N_samples116
48 must be 1 117
49 118
50 # initialize all actions of binary mask to 0 119
51 A_binary_mask = np.zeros((N_a,N_samples)) 120
52 # set the action that was chosen to 1 121
53 A_binary_mask[ind_of_action_taken,:] = 1 122
54 123
55 return A_binary_mask, ind_of_action_taken 124
56 125
57 126
58 127
59 # ===== activation functions and it's derivatives 128
60 ===== 129
61 # relu and its derivative
62 def relu(x):
63     return np.maximum(0,x)
64
65 def heaviside(x):
66     return np.heaviside(x,0)
67
68 # sigmoid and its derivative
69 def sigmoid(x):
70     return 1 / (1 + np.exp(-x))
71
72 def gradient_sigmoid(x):
73     return sigmoid(x) * (1 - sigmoid(x))
74
75 # tanh and its derivative
76 def tanh(x):
77     return np.tanh(x)
78
79 def gradient_tanh(x):
80     return 1 - np.tanh(x)**2
81
82 # identity and its derivative
83 def identity(x):
84     return x
85
86 def const(x):
87     return np.ones(x.shape)
88
89
90 def act_f_and_gradient(activation_function="relu"):150
91     if activation_function == "relu":
92         return relu, heaviside
93     elif activation_function == "sigmoid":
94         return sigmoid, gradient_sigmoid
95     elif activation_function == "tanh":
96         return tanh, gradient_tanh
97     else: # identity and constant 1
98         return identity, const
99
100
101
102 # ===== Replay Memory for Experience Replay (with 157
103 DQN) =====
104
105 Transition = namedtuple('Transition', ("state", " 158
106     action", "reward", "next_state", "done")) 159
107
108 class ReplayMemory(object):
109     def __init__(self, capacity):
110         self.memory = deque(maxlen=capacity)
111
112     def push(self, *args):
113         self.memory.append(Transition(*args))
114
115     def sample(self, batch_size):
116         # if less data than batch size, return all
117         data

```

```

163 # self.W2 = np.random.randn(self.O, self.K+1)/np.sqrt(self.K+1) # standard normal distribution, shape: (O, K+1)
164 # glorot/xavier normal initialization
165 self.W2 = np.random.standard_normal((self.O, self.K+1))*np.sqrt(2/(self.K+1 + self.O)) # standard normal distribution, shape: (O, K+1)
166 # self.W2 = np.random.randn(self.O, self.K+1) # standard normal distribution, shape: (O, K+1)
167
168 if self.method == "dqn":
169     self.W1_target = np.copy(self.W1)
170     self.W2_target = np.copy(self.W2)
171
172
173 def forward(self, x, target=False):
174     """
175     x has shape: (D+1, 1) (constant bias 1 must be added beforehand)
176     target: if True, use the weights of the target network
177
178     returns:
179         last logits (i.e. Qvalues) of shape (O, 1)
180
181     """
182     if target == True:
183         W1 = np.copy(self.W1_target)
184         W2 = np.copy(self.W2_target)
185     else:
186         W1 = np.copy(self.W1)
187         W2 = np.copy(self.W2)
188
189     # forward pass/propagation
190     a1 = W1 @ x
191     h1 = self.act_f_1(a1)
192     h1[0,:] = 1 # set first row (bias to second layer) to 1 (this ignores the weights for the k+1th hidden neuron, because this should not exist; this allows to only use matrix multiplication and simplify the gradients as we only need 2 instead of 4)
193     a2 = W2 @ h1
194     h2 = self.act_f_2(a2)
195     return a1, h1, a2, h2
196
197
198 def backward(self, R, x, Qvalues, Q_prime, a1, h1, a2, gamma, future_reward, action_binary_mask):
199     """
200     backward for methods "qlearning" and "sarsa"
201
202     x has shape (D+1, 1) (constant bias 1 must be added beforehand)
203     set future_reward=True for future reward with gamma>0, False for immediate reward
204     Q_prime must be chosen according to the method on x_prime (on- or off-policy)
205
206     """
207     # backward pass/backpropagation
208     # compute the gradient of the square loss with respect to the parameters
209
210     # ===== compute TD error (aka delta) =====
211     # make reward of shape (O, 1)
212     R_rep = np.tile(R, (self.O, 1))

```

```

275         self.dL_dW1 = ( (self.W2.T @ self.dL_da2) *
276             self.grad_act_f_1(a1) ) @ x.T
277
278
279     def update_parameters(self, eta):
280
281         # gradient clipping
282
283         # dL_dW1_norm = np.linalg.norm(self.dL_dW1)
284         # if dL_dW1_norm >= self.gradient_clip:
285             self.dL_dW1 = self.gradient_clip *
                self.dL_dW1 / dL_dW1_norm
286
287         # dL_dW2_norm = np.linalg.norm(self.dL_dW2)
288         # if dL_dW2_norm >= self.gradient_clip:
289             self.dL_dW2 = self.gradient_clip *
                self.dL_dW2 / dL_dW2_norm
290
291         # update W1 and W2
292         self.W2 = self.W2 + eta * self.dL_dW2
293         self.W1 = self.W1 + eta * self.dL_dW1
294
295
296
297
298     def train(self, env, N_episodes, eta, epsilon_0,
299         beta, gamma, alpha=0.001, gradient_clip=1,
300         batch_size=32, run_number=None):
301
302         """
303         alpha is used as weight for the exponential
304         moving average displayed during training
305
306         batch_size is only used for the DQN method.
307         """
308
309         # add training hyper parameters
310         self.N_episodes = N_episodes
311         self.eta = eta
312         self.epsilon_0 = epsilon_0
313         self.beta = beta
314         self.gamma = gamma
315         self.alpha = alpha
316         self.gradient_clip = gradient_clip
317         self.batch_size = batch_size
318
319
320         training_start = time.time()
321
322         try:
323
324             # initialize histories for important
325             metrics
326             self.R_history = np.full([self.
327                 N_episodes, 1], np.nan)
328             self.N_moves_history = np.full([self.
329                 N_episodes, 1], np.nan)
330             self.dL_dW1_norm_history = np.full([self.
331                 N_episodes, 1], np.nan)
332             self.dL_dW2_norm_history = np.full([self.
333                 N_episodes, 1], np.nan)
334
335             # progress bar
336             episodes = tqdm(np.arange(self.
337                 N_episodes), unit="episodes")
338             ema_previous = 0
339             n_steps = 0
340
341             for n in episodes:
342
343                 epsilon_f = self.epsilon_0 / (1 +
344                     beta * n) ## DECAYING EPSILON
345                 Done = 0

```

```

## SET DONE TO ZERO (BEGINNING
OF THE EPISODE)
i = 1

## COUNTER FOR NUMBER OF ACTIONS

S, X, allowed_a = env.
    Initialise_game() ##
INITIALISE GAME
X = np.expand_dims(X, axis=1)
## MAKE X A
TWO DIMENSIONAL ARRAY
X = np.copy(np.vstack((np.array
    ([[1]]), X))) # add bias term

if self.method == "sarsa":
    # compute Q values for the given
    state
    a1, h1, a2, Qvalues = self.
        forward(X) # -> shape (O,
        1)

    # choose an action A using
    epsilon-greedy policy
    A_binary_mask, A_ind =
        EpsilonGreedy_Policy(Qvalues
        , allowed_a, epsilon_f) #
        -> shape (O, 1)

while Done==0:

    ##
    START THE EPISODE

    if (self.method == "qlearning")
    or (self.method == "dqn"):
        # compute Q values for the
        given state
        a1, h1, a2, Qvalues = self.
            forward(X) # -> shape (
            O, 1)

        # choose an action A using
        epsilon-greedy policy
        A_binary_mask, A_ind =
            EpsilonGreedy_Policy(
            Qvalues, allowed_a,
            epsilon_f) # -> shape (
            O, 1)

        # take action and observe reward
        R and state S_prime
        S_prime, X_prime,
        allowed_a_prime, R, Done =
        env.OneStep(A_ind)
        X_prime = np.expand_dims(X_prime
        , axis=1)
        X_prime = np.copy(np.vstack((np.
        array([[1]]), X_prime))) #
        add bias term

        n_steps += 1

    if self.method == "dqn":

        # store the transition in
        memory
        self.replay_memory.push(X,
        A_ind, R, X_prime, Done)

        # sample a batch of
        transitions

```



```

372         transactions = self. 409
373         replay_memory.sample( 410
374             self.batch_size)
375         # turn list of transactions 411
376         into transaction of
377         lists 412
378         batch = Transition(*zip(*
379             transactions)) 413
380         # backward step and 414
381         parameter update 415
382         self.backward_dqn(batch,
383             self.gamma) 416
384         # update Q values indirectly by
385         updating the weights and
386         biases directly 417
387         if Done==1: # THE EPISODE HAS 418
388             ENDED, UPDATE...BE CAREFUL 419
389             THIS IS THE LAST STEP OF THE
390             EPISODE 420
391
392         if (self.method == "
393             qlearning") or (self.
394             method == "sarsa"):
395             # compute gradients and
396             update weights 421
397             self.backward(R, X, 422
398                 Qvalues, None, a1,
399                 h1, a2, None, 423
400                 future_reward=False,
401                 action_binary_mask= 424
402                 A_binary_mask) 425
403
404             # store history
405             # todo: record max possible
406             reward per episode 427
407             self.R_history[n] = np.copy(
408                 R) # reward per episode 428
409             self.N_moves_history[n] = np
410                 .copy(i) # nr moves per
411                 episode
412
413             # store norm of gradients
414             self.dL_dW1_norm_history[n] 429
415             = np.linalg.norm(self. 430
416                 dL_dW1) 431
417             self.dL_dW2_norm_history[n]
418             = np.linalg.norm(self. 432
419                 dL_dW2) 433
420
421             # compute exponential moving
422             average (EMA) to
423             display during training 435
424             ema = alpha*R + (1-alpha)* 436
425             ema_previous 437
426             if n == 0: # first episode 438
427                 ema = R
428             ema_previous = ema 439
429             if run_number is not None: 440
430                 episodes.set_description(
431                     (f"Run = {run_number
432                     }; EMA Reward = {ema
433                     :.2 f}") 443
434
435             else:
436                 episodes.set_description(
437                     (f"EMA Reward = {ema
438                     :.2 f}") 445
439
440             break 446
441
442         else: # IF THE EPISODE IS NOT
443             OVER... 448

```

```

if self.method == "qlearning":
    # chose next action off-
    policy
    Q_prime = np.max(self.
        forward(X_prime)
        [-1])

elif self.method == "sarsa":
    # chose next action on-
    policy

    a1_prime, h1_prime,
    a2_prime,
    Qvalues_prime = self
        .forward(X_prime) #
        -> shape (N_a, 1)

    # chose next action and
    save it
    A_binary_mask_prime,
    A_ind_prime =
        EpsilonGreedy_Policy
        (Qvalues_prime,
        allowed_a_prime,
        epsilon_f)

    # get Qvalue of next
    action
    Q_prime = Qvalues_prime[
        A_ind_prime]

if (self.method == "
    qlearning") or (self.
    method == "sarsa"):
    # backpropagation and
    weight update
    self.backward(R, X,
        Qvalues, Q_prime, a1
        , h1, a2, self.gamma
        , future_reward=True
        , action_binary_mask
        =A_binary_mask)

# NEXT STATE AND CO. BECOME
ACTUAL STATE...
if self.method == "sarsa":
    A_binary_mask = np.copy(
        A_binary_mask_prime)
    A_ind = np.copy(
        A_ind_prime)
    a1 = np.copy(a1_prime)
    h1 = np.copy(h1_prime)
    a2 = np.copy(a2_prime)
    Qvalues = np.copy(
        Qvalues_prime)
    S = np.copy(S_prime)
    X = np.copy(X_prime)
    allowed_a = np.copy(
        allowed_a_prime)

    i += 1 # UPDATE COUNTER FOR
    NUMBER OF ACTIONS

if (self.method == "dqn") and (
    n_steps % self.C == 0):
    # update target network
    every C steps
    self.W1_target = np.copy(
        self.W1)
    self.W2_target = np.copy(

```

```

449         self.W2)
450
451     training_end = time.time()
452     self.training_time_in_seconds =
453         training_end - training_start
454
455     return None
456
457 except KeyboardInterrupt as e:
458     # return nothing
459     training_end = time.time()
460     self.training_time_in_seconds =
461         training_end - training_start
462
463     return None
464
465 def save(self, name_extension=None):
466     # create directory for the model
467     name = f"{self.method}_{self.act_f_1_name}_{self.act_f_2_name}"
468     if name_extension is not None:
469         name += f"_{name_extension}"
470
471     path = f"models/{name}"
472     if not os.path.isdir(path): os.mkdir(path)
473     print(f"saving to: {path}")
474
475     # save weights
476     np.save(f"{path}/W1.npy", self.W1)
477     np.save(f"{path}/W2.npy", self.W2)
478
479     # save training history
480     np.save(f"{path}/training_history_R.npy", self.R_history)
481     np.save(f"{path}/training_history_N_moves.npy", self.N_moves_history)
482     np.save(f"{path}/training_history_dL_dW1_norm.npy", self.dL_dW1_norm_history)
483     np.save(f"{path}/training_history_dL_dW2_norm.npy", self.dL_dW2_norm_history)
484
485     # save training parameters and other general info
486     params = {
487         "method": self.method,
488         "N_episodes": self.N_episodes,
489         "eta": self.eta,
490         "epsilon_0": self.epsilon_0,
491         "beta": self.beta,
492         "gamma": self.gamma,
493         "alpha": self.alpha,
494         # "gradient_clip": self.gradient_clip,
495         "seed": self.seed,
496         "D": self.D,
497         "K": self.K,
498         "O": self.O,
499         "training_time_in_seconds": self.training_time_in_seconds
500     }
501     if self.method == "dqn":
502         params["capacity"] = self.capacity
503         params["batch_size"] = self.batch_size
504         params["C"] = self.C
505     with open(f"{path}/training_parameters.json", "w") as f:
506         json.dump(params, f)
507
508     def load_from(method, act_f_1, act_f_2,
509
510         name_extension=None):
511
512     # read values and store in neural network instance
513     name = f"{method}_{act_f_1}_{act_f_2}"
514     if name_extension is not None:
515         name += f"_{name_extension}"
516
517     path = f"models/{name}"
518     # print(f"loading from: {path}")
519
520     # initialize neural network
521     nn = NeuralNetwork(0,0,0, activation_function_1=
522         act_f_1, activation_function_2=act_f_2,
523         method=method)
524
525     # network weights
526     nn.W1 = np.load(f"{path}/W1.npy")
527     nn.W2 = np.load(f"{path}/W2.npy")
528
529     # network training history
530     nn.R_history = np.load(f"{path}/training_history_R.npy")
531     nn.N_moves_history = np.load(f"{path}/training_history_N_moves.npy")
532     nn.dL_dW1_norm_history = np.load(f"{path}/training_history_dL_dW1_norm.npy")
533     nn.dL_dW2_norm_history = np.load(f"{path}/training_history_dL_dW2_norm.npy")
534
535     # network training parameters
536     with open(f"{path}/training_parameters.json", "r") as f:
537         params = json.load(f)
538
539     # set parameters to the network instance
540     nn.method = params["method"]
541     nn.N_episodes = int(params["N_episodes"])
542     nn.eta = float(params["eta"])
543     nn.epsilon_0 = float(params["epsilon_0"])
544     nn.beta = float(params["beta"])
545     nn.gamma = float(params["gamma"])
546     nn.alpha = float(params["alpha"])
547     # nn.gradient_clip = float(params["gradient_clip"])
548
549     try:
550         nn.seed = int(params["seed"])
551     except:
552         nn.seed = params["seed"]
553     nn.D = int(params["D"])
554     nn.K = int(params["K"])
555     nn.O = int(params["O"])
556     nn.training_time_in_seconds = float(params["training_time_in_seconds"])
557
558     if nn.method == "dqn":
559         nn.capacity = int(params["capacity"])
560         nn.batch_size = int(params["batch_size"])
561         nn.C = int(params["C"])
562
563     if nn.method == "dqn":
564         nn.W1_target = np.copy(nn.W1)
565         nn.W2_target = np.copy(nn.W2)
566
567     return nn

```

Listing 1: caption